

# The Nile Programming Language: Declarative Stream Processing for Media Applications

Dan Amelang  
Viewpoints Research Institute  
UCSD PhD Student

*Southern California Programming Languages and  
Systems Workshop – Spring 2011*

# 2D Vector Graphics

The image shows two overlapping windows. The top window is OpenOffice Writer, displaying a document with the text "The M... John, by the grace of God King of Engla... and Count of Anjou, to his archbishops, bishops stewards, servants, and to all his officials and lo...". The bottom window is Mozilla Firefox Start Page, showing the Google search bar and various links. Overlaid on the Firefox window are two diagrams illustrating vector graphics for the characters 'H' and 'f'. The 'H' diagram shows a black box with annotations for "Left to right baseline origin", "Right side bearing", "Advance width", and "Left side bearing". The 'f' diagram shows a black box with annotations for "Right to left baseline origin", "Upright baseline offset", "Right side bearing", "Left side bearing", and "Advance width".

## Designing an Integrated Review Sheet for an Electronic Textbook

Neema Moraveji, Abigail Travis, Maura Bidinost, and Matthew Halpern  
 neema@cmu.com, atravis@cs.cmu.edu, maura@cmu.edu, matt@thesindicate.org  
 Human Computer Interaction Institute  
 Carnegie Mellon University  
 Pittsburgh, PA 15213 USA

### ABSTRACT

In this paper, we present findings and design decisions arisen while designing a review sheet within the confines of a pre-existing digital textbook, AdaptiveBook. Through user studies, we found that instructors and students cited a lack of both context and integration as the major problems in current high-tech teaching and studying tools. Because research on electronic books indicates that metaphors

obstacles. From these user studies, we formed the concept of the integrated review sheet.

### AdaptiveBook

We were presented with AdaptiveBook as a platform for the review sheet. AdaptiveBook is textbook reader software that enables users to create 'markups' on the textbook including highlights, annotations, and general notes. These markups can thereafter be copied and shared

The image shows the JellyfishSSH 4.2 interface. It features a "Group Filter" set to "Show All" and a "Bookmark" dropdown set to "Select...". Below these are two tabs: "Connection Settings" and "SSH Options". The "Connection Settings" tab is active, showing fields for "Bookmark Title" (CASLab Zeus), "Host or IP" (zeus.caslab.queensu.ca), "Port" (22), "Login name" (YOUR\_ID), "Shell" (/bin/bash), "Type" (SSH 2), and "Member of Group" (Main). At the bottom are buttons for "Add", "Delete", "Clear All", and "Connect".

of providing this context usually fell on the student to

# State of the Art – 10,000+ lines of...

Cairo (Firefox, Linux)

```
...
if (lxi > fill_start)
{
    ADD_SATURATE_8 (ap + fill_start,
                    fill_size * N_X_FRAC (8),
                    lxi - fill_start);
    fill_start = lxi;
}
else if (lxi < fill_start)
{
    ADD_SATURATE_8 (ap + lxi,
                    N_X_FRAC (8),
                    fill_start - lxi);
}
if (rxix < fill_end)
{
    ADD_SATURATE_8 (ap + rxix,
                    fill_size * N_X_FRAC (8),
                    fill_end - rxix);
}
...
```

Skia (Chrome, Android)

```
...
if (currE->fLastY == curr_y) {
    if (currE->fCurveCount < 0) {
        if (((SkCubicEdge*)currE)->updateCubic())
        {
            SkASSERT(currE->fFirstY == curr_y + 1);
            newX = currE->fX;
            goto NEXT_X;
        }
    } else if (currE->fCurveCount > 0) {
        if (((SkQuadraticEdge*)currE)->updateQuadratic())
        {
            newX = currE->fX;
            goto NEXT_X;
        }
    }
}
...
```

```

type Color      = (a, r, g, b : Real)
type Point     = (x, y : Real)
type Matrix    = (a, b, c, d, e, f : Real)
type Bezier    = (A, B, C : Point)
type EdgeSpan  = (x, y, c, l : Real)
type EdgeSample = (x, y, a, h : Real)
type PixelCoverage = (x, y, c, ic : Real)

type Texture   = Point >> Color
type Compositor = (Color, Color) >>> Color

| (a : Real) | : Real
{ -a if a < 0, a }

(a : Real) < (b : Real) : Real
{a if a < b, b}

(a : Real) > (b : Real) : Real
{a if a > b, b}

(a : Real) ~ (b : Real) : Real
(a + b) / 2

(M : Matrix) ⊗ (A : Point) : Point
(M.a × A.x + M.c × A.y + M.e, M.b × A.x + M.d × A.y + M.f)

TransformBeziers (M : Matrix) : Bezier >> Bezier
∀ (A, B, C)
  >> (M ⊗ A, M ⊗ B, M ⊗ C)

UniformColor (C : Color) : Texture
∀ _
  >> (C.a, C.a × C.r, C.a × C.g, C.a × C.b)

CompositeOver : Compositor
∀ (a, b)
  >> a + b × (1 - a.a)

ClipBeziers (min, max : Point) : Bezier >> Bezier
∀ (A, B, C)
  bmin = A < B < C
  bmax = A > B > C
  inside = min ≤ bmin ∧ bmax ≤ max
  outside = bmax ≤ min ∨ max ≤ bmin
  if inside.x ∧ inside.y
    >> (A, B, C)
  else if outside.x ∨ outside.y
    cA = min > A < max
    cC = min > C < max
    >> (cA, cA ~ cC, cC)
  else
    ABBC = (A ~ B) ~ (B ~ C)
    nearmin = | ABBC - min | < 0.1
    nearmax = | ABBC - max | < 0.1
    M = {min if nearmin, max if nearmax, ABBC}
    << (M, B ~ C, C) << (A, A ~ B, M)

```

```

CompositeTextures (t1 : Texture, t2 : Texture, c : Compositor) : Texture
  ⇒ DupZip (t1, t2) → c

ExpandSpans : EdgeSpan >> PixelCoverage
  ∀ (x, y, c, l)
    if c
      >> (x, y, c, 1 - c)
    if l > 0
      << (x + 1, y, 1, l - 1)

ExtractSamplePoints : PixelCoverage >> Point
  ∀ (x, y, _, _)
    >> (x, y)

ApplyTexture (t : Texture) : EdgeSpan >> (Color, PixelCoverage)
  ⇒ ExpandSpans → DupZip (ExtractSamplePoints → t, (-))

DecomposeBeziers : Bezier >> EdgeSample
  ∀ (A, B, C)
    inside = (| A | = | C | ∨ | A | = | C |)
    if inside.x ∧ inside.y
      P = | A | < | C |
      w = P.x + 1 - (C.x ~ A.x)
      h = C.y - A.y
      >> (P.x + 1/2, P.y + 1/2, w × h, h)
    else
      ABBC = (A ~ B) ~ (B ~ C)
      min = | ABBC |
      max = | ABBC |
      nearmin = | ABBC - min | < 0.1
      nearmax = | ABBC - max | < 0.1
      M = {min if nearmin, max if nearmax, ABBC}
      << (M, B ~ C, C) << (A, A ~ B, M)

CombineEdgeSamples : EdgeSample >> EdgeSpan
  (x, y, A, H) = 0
  ∀ (x', y', a, h)
    if y' = y
      if x' = x
        A' = A + a
        H' = H + h
      else
        l = {x' - x - 1 if |H| > 0.5, 0}
        >> (x, y, |A| < 1, l)
        A' = H + a
        H' = H + h
    else
      >> (x, y, |A| < 1, 0)
      A' = a
      H' = h
  >> (x, y, |A| < 1, 0)

```

```

Rasterize : Bezier >> EdgeSpan
  ⇒ DecomposeBeziers → SortBy (@x) → SortBy (@y) → CombineEdgeSamples

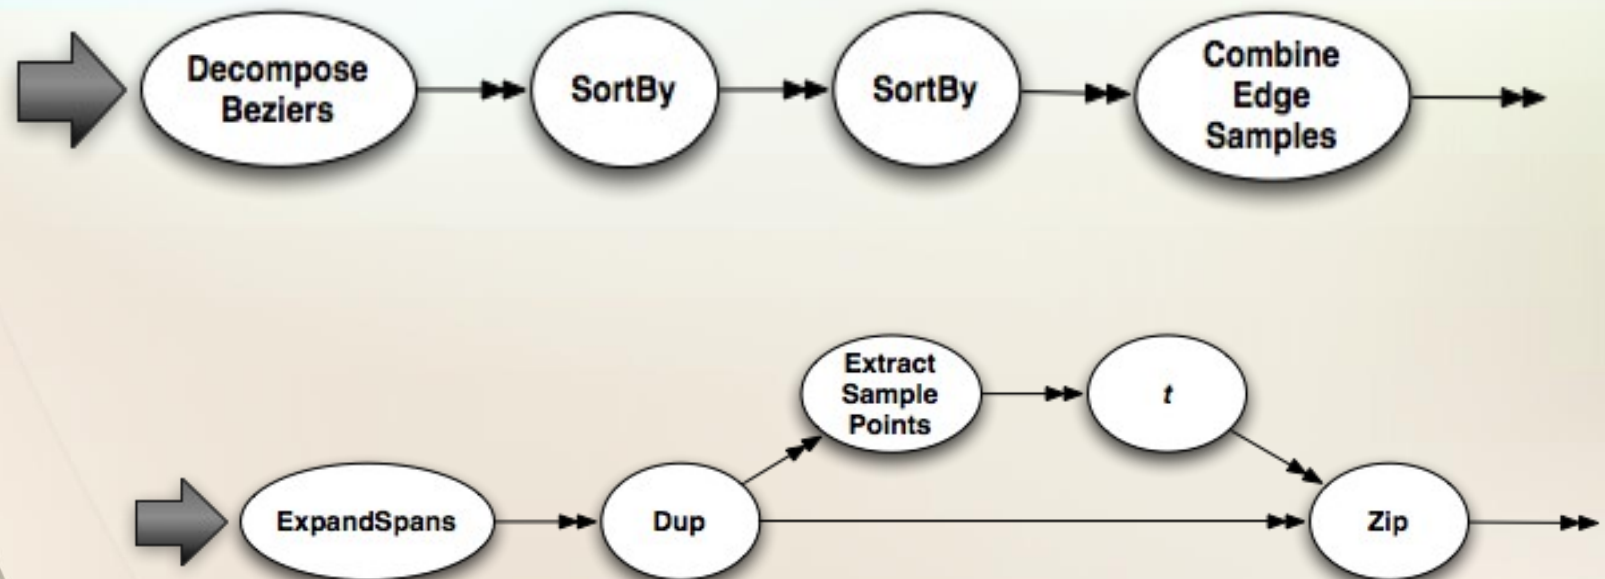
```

# Talk Overview


- The Nile language
  - Dataflow
  - Data manipulation
- Graphics rendering in Nile
- Parallelism in Nile
- The “Frankenstein” Application
- Future Work
- Related Work

# Nile Dataflow

- Processes communicate asynchronously via unidirectional, homogenous, typed streams of structured data
- Single input / single output (with some exceptions)



# Rendering Pipeline




TransformBeziers (matrix) →  
ClipBeziers (min, max) →  
Rasterize →  
ApplyTexture (texture) →  
WriteToImage (image)



# Rasterization Pipeline

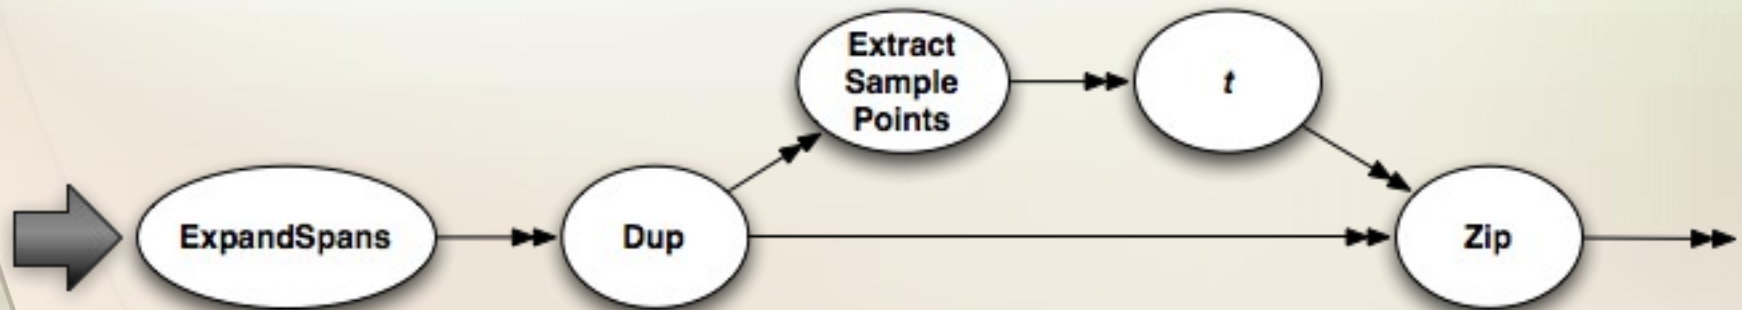
**Rasterize : Bezier >> EdgeSpan**  
⇒ **DecomposeBeziers** →  
**SortBy (@x)** → **SortBy (@y)** →  
**CombineEdgeSamples**





# Texturing Pipeline

```
ApplyTexture (t : Texture) : EdgeSpan >> (Color, PixelCoverage)  
  ⇒ ExpandSpans →  
    DupZip (ExtractSamplePoints → t, (-))
```



# Data Manipulation

- Statically typed with type inference
- Single assignment
- User-defined record types and operators
- Pattern matching
- Syntax for stream I/O

# Data Manipulation

**Sum : Real >> Real**

**s = 0**

**∀ x**

**s' = s + x**

**>> s**

**>> s**

# User-defined Types and Operators

```
type Point = (x, y : Real)  
type Matrix = (a, b, c, d, e, f : Real)  
type Bezier = (A, B, C : Point)
```

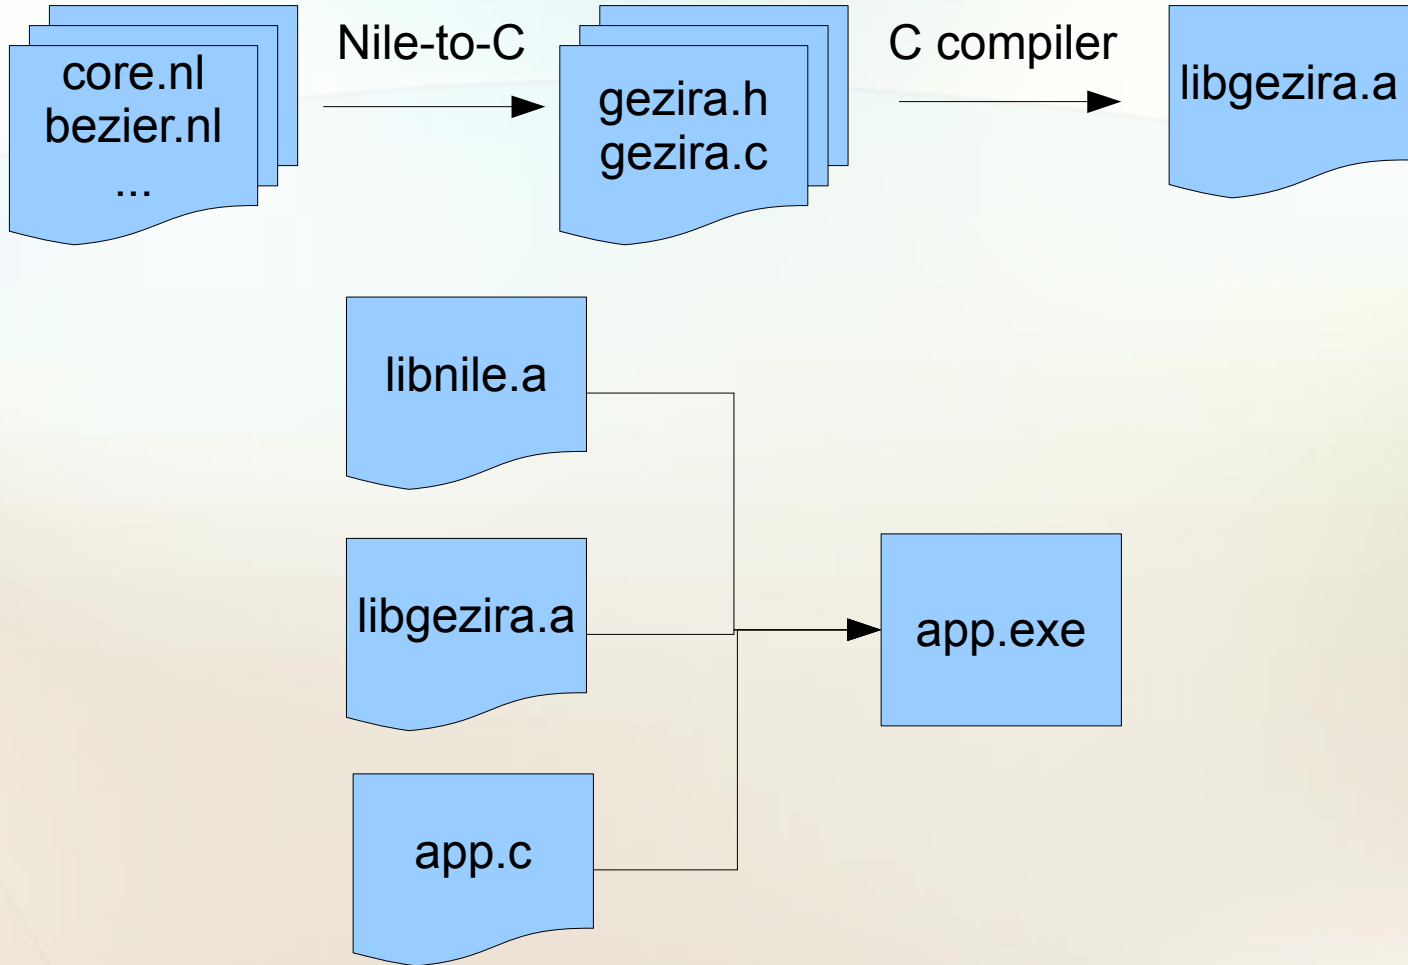
```
(M : Matrix) × (A : Point) : Point  
(M.a × A.x + M.c × A.y + M.e,  
  M.b × A.x + M.d × A.y + M.f)
```

```
TransformBeziers (M : Matrix) : Bezier >> Bezier  
  V (A, B, C)  
    >> (M × A, M × B, M × C)
```

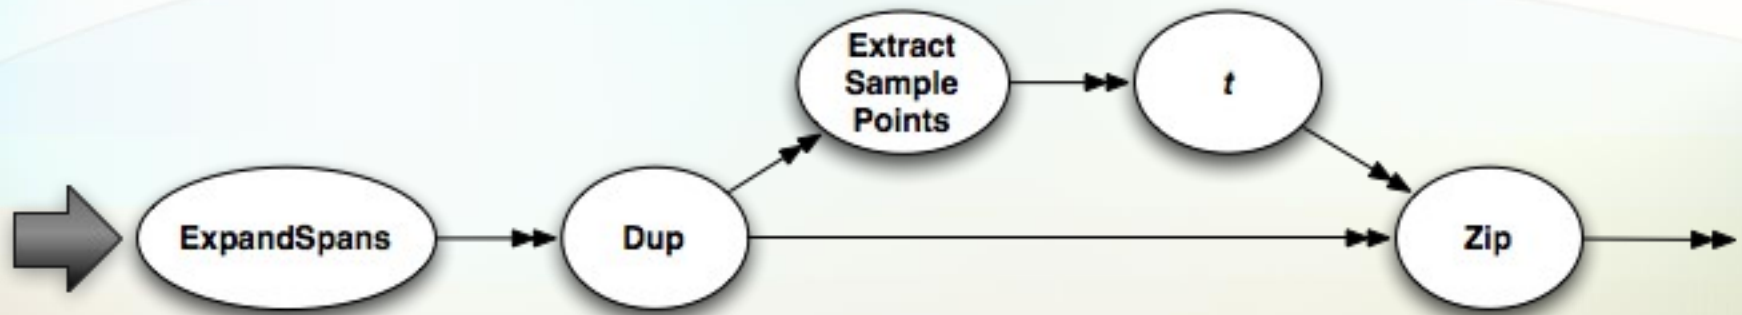
# Gezira: A 2D Vector Graphics Renderer Written in Nile

- In under 400 lines of code:
  - Anti-aliased rasterization of Bezier shapes
  - Affine transformation
  - Geometry clipping
  - 26 compositing operators
  - Texture transformation and extend styles
  - Bilinear and bicubic filters
  - Gaussian blur
  - Multi-stop linear and radial color gradients
  - Pen stroke paths with 3 join styles and 3 cap styles
  - Geometry bounds calculation

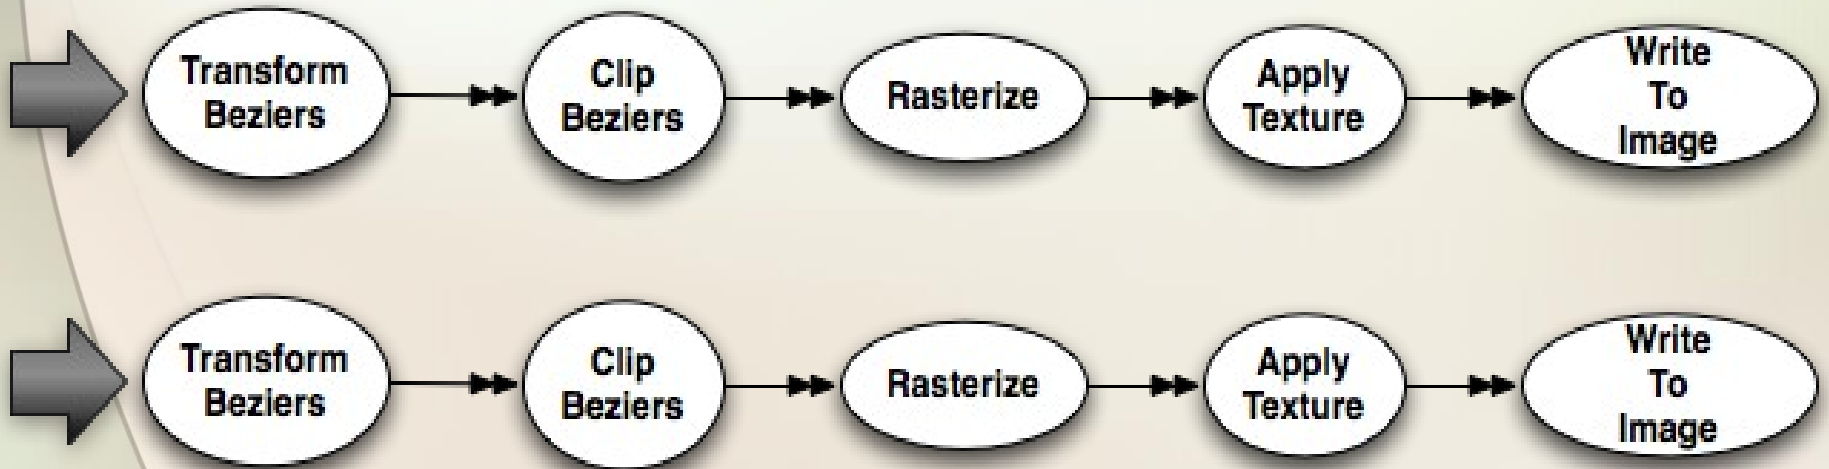

# Workflow



# Intra-pipeline Parallelism



# Inter-pipeline Parallelism



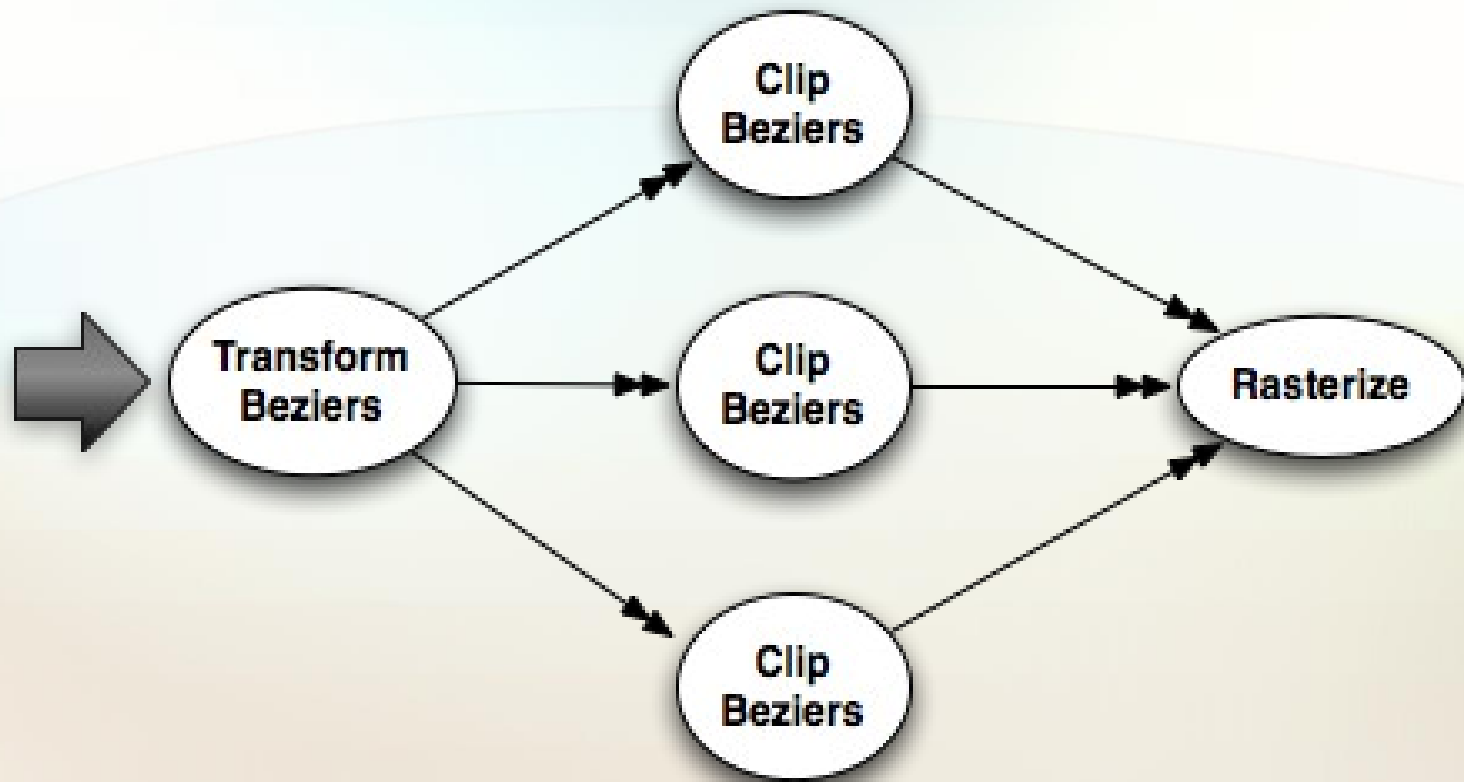


# Instruction-level SIMD Parallelism

```
(M : Matrix) × (A : Point) : Point  
  (M.a × A.x + M.c × A.y + M.e,  
   M.b × A.x + M.d × A.y + M.f)
```

```
TransformBeziers (M : Matrix) : Bezier >> Bezier  
  ∀ (A, B, C)  
    >> (M × A, M × B, M × C)
```


# Process-level Data Parallelism



# Nile Runtime

- Multi-threaded
- Load balancing
- Heap balancing
- User space synchronization

# Speedup on 40 core machine



# Meet “Frank”



# Future Work

- More applications
  - Data compression (zlib, png)
  - Audio/video decoding
  - 3D graphics
  - Image processing
- More compiler backends
  - OpenCL (for GPUs)
  - Javascript
- More language features
  - Sliding windows
  - Feedback networks
  - Binary streams

# Related Work

- Programming models
  - Kahn Process Networks [Kahn74, Kahn76]
  - Dataflow Process Networks [Lee95]
- Programming languages
  - Lucid [Wadge85]
  - VAL [Acherman79], later SISAL [Mcgraw85]
  - Id [Arvind90], later Parallel Haskell (pH) [Nikhil93]
  - Streamit [Thies02]
  - Single Assignment C [Scholz03] and S-Net [Grelck07]





**Extra slides...**

# Model of Graphics Rendering?

From Computer Graphics: Principles and Practice [Foley et al.]:

Once the ET has been formed, the following processing steps for the scan-line algorithm are completed:

1. Set  $y$  to the smallest  $y$  coordinate that has an entry in the ET, that is,  $y$  for the first nonempty bucket.
2. Initialize the AET to be empty.
3. Repeat until the AET and ET are empty;
  - (a) Move from ET bucket  $y$  to the AET edges whose  $y_{\min} = y$  (add entering edges).
  - (b) Remove from the AET those entries for which  $y = y_{\max}$  (edges not involved in the next scan line), then sort the AET on  $x$ .
  - (c) Fill in desired pixel values on scan line  $y$  by using pairs of  $x$  coordinates from the AET (suitably rounded).
  - (d) Increment  $y$  by 1 (to the coordinate of the next scan line).
  - (e) For each nonvertical edge remaining in the AET, update  $x$  for the new  $y$ .

Given the x and y coordinates of the lower-left corner of a pixel, the coverage contribution of an edge  $\overrightarrow{AB}$  can be calculated as follows:

$$\sigma(P, Q) = (Q_y - P_y)(x + 1 - \frac{Q_x + P_x}{2})$$

$$\gamma(P) = \begin{aligned} & \min(x + 1, \max(x, P_x)), \\ & \min(y + 1, \max(y, P_y)) \end{aligned}$$

$$\omega(P) = \begin{aligned} & \frac{1}{m}(\gamma(P)_y - P_y) + P_x, \\ & m(\gamma(P)_x - P_x) + P_y \end{aligned}$$

$$\begin{aligned} \text{coverage}(\overrightarrow{AB}) &= \sigma(\gamma(A), \gamma(\omega(A))) + \\ & \sigma(\gamma(\omega(A)), \gamma(\omega(B))) + \\ & \sigma(\gamma(\omega(B)), \gamma(B)) \end{aligned}$$

The total coverage contribution of a polygon is the linear combination of the edge contributions, with some additional adjustment:

$$\min(|\sum coverage(\overrightarrow{AB}_i)|, 1)$$

**TODO: size graph of cairo vs. gezira**